Natural Language Processing
MEMMs and CRFs

Felipe Bravo-Marquez

October 21, 2019

MEMMs

- The goal of sequence labeling is to assign tags to words, or more generally, to assign discrete labels to discrete elements in a sequence [Eisenstein, 2018].
- Well known examples of this problem are: part-of-speech tagging (POS) and Named Entity Recognition (NER).
- Maximum-entropy Markov models (MEMMs) make use of log-linear multi-class models for sequence labeling tasks [McCallum et al., 2000].
- In the early NLP literature, logistic regression was often called maximum entropy classification [Eisenstein, 2018].
- Hence, MEMMs will look very similar to the multi-class softmax models seen in the lecture about linear models.
- In contrast to HMMs, here we rely on parameterized functions.

MEMMs

- The goal of MEMMs is model the following conditional distribution:

$$P(s_1, s_2 \ldots, s_m | x_1, \ldots, x_m)$$

- Where each x_j for $j = 1 \ldots m$ is the j-th input symbol (for example the j-th word in a sentence), and each s_j for $j = 1 \ldots m$ is the j-th tag.[1]

- We would expect P(DET,NOUN,VERB|the,dog,barks) to be higher than P(VERB,VERB,VERB|the,dog,barks) in a model trained from training dataset.

---

[1] These slides are based on lecture notes of Michael Collins
http://www.cs.columbia.edu/~mcollins/crf.pdf. The notation and terminology has been adapted to be consistent with the rest of the course.

MEMMs

- We use S to denote the set of possible tags.
- We assume that S is a finite set.
- For example, in part-of-speech tagging of English, S would be the set of all possible parts of speech in English (noun, verb, determiner, preposition, etc.).
- Given a sequence of words $x_1, \ldots, x_m$, there are k^m possible part-of-speech sequences $s_1, \ldots, s_m$, where $k = |S|$ is the number of possible parts of speech.
- We want to estimate a distribution over these k^m possible sequences.

MEMMs

- In a first step, MEMMs use the following decomposition (s_0 has always a special tag $*$):

$$
\begin{aligned}
P(s_1, s_2 \ldots, s_m | x_1, \ldots, x_m) &= \prod_{i=1}^{m} P(s_i | s_1 \ldots, s_{i-1}, x_1, \ldots, x_m) \\
&= \prod_{i=1}^{m} P(s_i | s_{i-1}, x_1, \ldots, x_m)
\end{aligned}
\tag{1}
$$

- The first equality is exact (it follows by the chain rule of conditional probabilities).
- The second equality follows from an independence assumption, namely that for all i,

$$
P(s_i | s_1 \ldots, s_{i-1}, x_1, \ldots, x_m) = P(s_i | s_{i-1}, x_1, \ldots, x_m)
$$

MEMMs

- Hence we are making a first order Markov assumption similar to the Markov assumption made in HMMs[2].
- The tag in the i-th position depends only on the tag in the $(i-1)$-th position.
- Having made these independence assumptions, we then model each term using a multiclass log-linear (softmax) model:

$$P(s_i|s_{i-1}, x_1, \ldots, x_m) = \frac{\exp(\vec{w} \cdot \vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s_i))}{\sum_{s' \in S} \exp(\vec{w} \cdot \vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s'))} \quad (2)$$

---

[2]We actually made a second order Markov assumption in HMMs. MEMMs can also be extended to second order assumptions.

MEMMs

Here $\vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s_i)$ is a feature vector where:

- $x_1, \cdots, x_m$ is the entire sentence being tagged.
- i is the position to be tagged (can take any value from 1 to m).
- s_{i-1} is the previous tag value (can take any value in S).
- s_i is the new tag value (can take any value in S).

Example of Features used in Part-of-Speech Tagging

1. $\vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s_i)_{[1]} = 1$ if s_i = ADVERB and word x_i ends in "-ly"; 0 otherwise.
   If the weight $\vec{w}_{[1]}$ associated with this feature is large and positive, then this feature is essentially saying that we prefer labelings where words ending in -ly get labeled as ADVERB.

2. $\vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s_i)_{[2]} = 1$ if $i = 1$, s_i = VERB, and x_m =?; 0 otherwise.
   If the weight $\vec{w}_{[2]}$ associated with this feature is large and positive, then labelings that assign VERB to the first word in a question (e.g., "Is this a sentence beginning with a verb?") are preferred.

3. $\vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s_i)_{[3]} = 1$ if s_{i-1} = ADJECTIVE and s_i = NOUN; 0 otherwise.
   Again, a positive weight for this feature means that adjectives tend to be followed by nouns.

4. $\vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s_i)_{[4]} = 1$ if s_{i-1} = PREPOSITION and s_i = PREPOSITION.
   A negative weight $\vec{w}_{[4]}$ for this function would mean that prepositions don't tend to follow prepositions.

[3]Source: https://blog.echen.me/2012/01/03/
introduction-to-conditional-random-fields/

Feature Templates

It is possible to define more general feature templates covering unigrams, bigrams, n-grams of words as well as tag values of s_{i-1} and s_i.

1. A word unigram and tag unigram feature template:
   $\vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s_i)_{[index(j,z)]} = 1$ if $s_i = \text{TAG}_{[j]}$ and $x_i = \text{WORD}_{[z]}$; 0 otherwise $\forall j, z$.
   Notice that j is and index spanning all possible tags in S and z is another index spanning the words in the vocabulary V.

2. A word bigram and tag bigram feature template:
   $\vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s_i)_{[index(j,z,u,v)]} = 1$ if $s_{i-1} = \text{TAG}_{[j]}$ and $s_i = \text{TAG}_{[z]}$ and $x_{i-1} = \text{WORD}_{[u]}$ and $x_i = \text{WORD}_{[v]}$; 0 otherwise $\forall j, z, u, v$.

The function $index(j, k, ...)$ will map each different feature to a unique index in the feature vector.

Notice that the resulting vector will be very high-dimensional and sparse.

MEMMs and Multi-class Softmax

- Notice that the log-linear model from above is very similar to the multi-class softmax model presented in the lecture about linear models.

- A general log-linear model has the following form:

$$P(y|x; \vec{w}) = \frac{\exp(\vec{w} \cdot \vec{\phi}(x, y))}{\sum_{y' \in Y} \exp(\vec{w} \cdot \vec{\phi}(x, y'))}$$

- A multi-class softmax model has the following form:

$$\begin{aligned}
\hat{\vec{y}} &= \text{softmax}(\vec{x} \cdot W + \vec{b}) \\
\hat{\vec{y}}_{[i]} &= \frac{e^{(\vec{x} \cdot W + \vec{b})_{[i]}}}{\sum_j e^{(\vec{x} \cdot W + \vec{b})_{[j]}}}
\end{aligned} \tag{3}$$

MEMMs and Multi-class Softmax

- Difference 1: in the log-linear model we have a fixed parameter vector $\vec{w}$ instead of having multiple vectors (one column of W for each class value).

- Difference 2: the feature vector of the log-linear model $\vec{\phi}(x, y)$ includes information of the label y, whereas the input vector $\vec{x}$ of the softmax model is independent of y.

- Log-linear models allow using features that consider the interaction between x and y (e.g., x ends in "ly" and y is an ADVERB).

Training MEMMs

- Once we've defined the feature vectors $\vec{\phi}$, we can train the parameters $\vec{w}$ of the model in the usual way linear models are trained.
- We set the negative log-likelihood as the loss function and optimize parameters using gradient descent from the training examples.
- This is equivalent as using the cross-entropy loss.
- "Any loss consisting of a negative log-likelihood is a cross-entropy between the empirical distribution defined by the training set and the probability distribution defined by model" [Goodfellow et al., 2016].

Decoding with MEMMs

- The decoding problem is as follows.
- We are given a new test sequence $x_1, \ldots, x_m$.
- Our goal is to compute the most likely state sequence for this test sequence,

$$\arg \max_{s_1, \ldots, s_m} P(s_1, \ldots, s_m | x_1, \ldots, x_m). \qquad (4)$$

- There are k^m possible state sequences, so for any reasonably large sentence length m brute-force search through all the possibilities will not be possible.
- We can use the Viterbi alogrithm in a similar way as used for HMMs.

Decoding with MEMMs

- The basic data structure in the algorithm will be a dynamic programming table π with entries $\pi[j, s]$ for $j = 1, \ldots, m$, and $s \in S$.
- $\pi[j, s]$ will store the maximum probability for any state sequence ending in state s at position j.
- More formally, our algorithm will compute

$$\pi[j, s] = \max_{s_1, \ldots, s_{j-1}} \left(P(s|s_{j-1}, x_1, \ldots, x_m) \prod_{k=1}^{j-1} P(s_k|s_{k-1}, x_1, \ldots, x_m) \right)$$

for all $j = 1, \ldots, m$, and for all $s \in S$.

Decoding with MEMMs

The algorithm is as follows:

- Initialization: for $s \in S$

$$\pi[1, s] = P(s|s_0, x_1, \ldots, x_m)$$

where s_0 is a special "initial" state.

- For $j \in \{2, \ldots, m\}$, $s \in \{1, \ldots, k\}$

$$\pi[j, s] = \max_{s' \in S} \left[\pi[j-1, s'] \times P(s|s', x_1, \ldots, x_m) \right]$$

Decoding with MEMMs

- Finally, having filled in the $\pi[j, s]$ values for all j, s, we can calculate

$$\max_{s_1,\ldots,s_m} = \max_s \pi[m, s].$$

- The algorithm runs in $O(mk^2)$ time (i.e., linear in the sequence length m, and quadratic in the number of tags k).

- As in the Viterbi algorithm for HMMs, we can compute the highest-scoring sequence using backpointers in the dynamic programming algorithm.

Comparison between MEMMs and HMMs

- So what is the motivation for using MEMMs instead of HMMs?
- Note that the Viterbi decoding algorithms for the two models are very similar.
- In MEMMs, the probability associated with each state transition s_{i-1} to s_i is

$$P(s_i|s_{i-1}, x_1, \ldots, x_m) = \frac{\exp(\vec{w} \cdot \vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s_i))}{\sum_{s' \in S} \exp(\vec{w} \cdot \vec{\phi}(x_1, \cdots, x_m, i, s_{i-1}, s'))}$$

- In HMMs, the probability associated with each transition is:

$$P(s_i|s_{i-1}, x_1, \ldots, x_m) = P(s_1|s_{i-1})P(x_i|s_i)$$

Comparison between MEMMs and HMMs

- The key advantage of MEMMs is that the use of feature vectors $\vec{\phi}$ allows much richer representations than those used in HMMs.
- For example, the transition probability can be sensitive to any word in the input sequence $x_1, \ldots, x_m$.
- In addition, it is very easy to introduce features that are sensitive to spelling features (e.g., prefixes or suffixes) of the current word x_i, or of the surrounding words.
- These features are useful in many NLP applications, and are difficult to incorporate within HMMs in a clean way.

Conditional Random Fields (CRFs)

- We now turn to Conditional Random Fields (CRFs) [Lafferty et al., 2001].
- Notation: for convenience, we'll use $x_{1:m}$ to refer to an input sequence $x_1, \ldots, x_m$, and $s_{1:m}$ to refer to a sequence of tags $s_1, \ldots, s_m$.
- The set of all possible tags is again S.
- The set of all possible tag sequences is S^m.
- In conditional random fields we'll again build a model of

$$P(s_1, \ldots, s_m | x_1, \ldots, x_m) = P(s_{1:m} | x_{1:m})$$

Conditional Random Fields (CRFs)

- A first key idea in CRFs will be to define a feature vector that maps an entire input sequence $x_{1:m}$ paired with an entire tag sequence $s_{1:m}$ to some d-dimensional feature vector:

$$\vec{\Phi}(x_{1:m}, s_{1:m}) \in \mathcal{R}^d$$

- We'll soon give a concrete definition for $\vec{\Phi}$.
- For now just assume that some definition exists.
- We will often refer to $\vec{\Phi}$ as being a "global" feature vector.
- It is global in the sense that it takes the entire state sequence into account.

Conditional Random Fields (CRFs)

- In CRFs we build a giant log-linear model:

$$P(s_{1:m}|x_{1:m}; \vec{w}) = \frac{\exp(\vec{w} \cdot \vec{\Phi}(x_{1:m}, s_{1:m}))}{\sum_{s'_{1:m} \in S^m} \exp(\vec{w} \cdot \vec{\Phi}(x_{1:m}, s'_{1:m}))}$$

- This is "just" another log-linear model, but it is is "giant".
- The space of possible values for $s_{1:m}$ is huge S^m.
- The normalization constant (denominator in the above expression) involves a sum over all possible tag sequences S^m.
- These issues might seem to cause severe computational problems.
- Under appropriate assumptions we can train and decode efficiently with this type of model.

Conditional Random Fields (CRFs)

- We define the global feature vector $\vec{\Phi}(x_{1:m}, s_{1:m})$ as follows:

$$\vec{\Phi}(x_{1:m}, s_{1:m}) = \sum_{j=1}^{m} \vec{\phi}(x_{1:m}, j, s_{j-1}, s_j)$$

  where $\vec{\phi}(x_{1:m}, j, s_{j-1}, s_j)$ are the same as the feature vectors used in MEMMs.

- Example:
  $\vec{\Phi}([\text{the,dog,barks}], \text{DET,NOUN,VERB}]) = \vec{\phi}([\text{the,dog,barks}], 1, *, \text{DET}) + \vec{\phi}([\text{the,dog,barks}], 2, \text{DET, NOUN}) + \vec{\phi}([\text{the,dog,barks}], 3, \text{NOUN, VERB})$

- Essentially, we are adding up many sparse vectors.

Conditional Random Fields (CRFs)

- We are assuming that for any dimension of $\vec{\Phi}_{[k]}, k = 1, \ldots, d$, the k'th global feature is:

$$\vec{\Phi}(x_{1:m}, s_{1:m})_{[k]} = \sum_{j=1}^{m} \vec{\phi}(x_{1:m}, j, s_{j-1}, s_j)_{[k]}$$

- Thus $\vec{\Phi}(x_{1:m}, s_{1:m})_{[k]}$ is calculated by summing the "local" feature vector $\vec{\phi}(x_{1:m}, j, s_{j-1}, s_j)_{[k]}$ over the m different tag transitions in $s_1, \ldots, s_m$.

- We would expect each local vector to encode relevant information about the tag transition by turning on some vector dimensions (setting the value to one).

- We now turn to two critical practical issues in CRFs: first, decoding, and second, parameter estimation (training).

Decoding with CRFs

- The decoding problem in CRFs is as follows.
- For a given input sequence $x_{1:m} = x_1, x_2, \ldots, x_m$, we would like to find the most likely underlying state sequence under the model, that is,

$$
\begin{aligned}
arg \max_{s_{1:m} \in S^m} P(s_{1:m}|x_{1:m}; \vec{w}) \quad &= arg \max_{s_{1:m} \in S^m} \frac{\exp(\vec{w} \cdot \vec{\Phi}(x_{1:m}, s_{1:m}))}{\sum_{s'_{1:m} \in S^m} \exp(\vec{w} \cdot \vec{\Phi}(x_{1:m}, s'_{1:m}))} \\
&= arg \max_{s_{1:m} \in S^m} \exp(\vec{w} \cdot \vec{\Phi}(x_{1:m}, s_{1:m})) \\
&= arg \max_{s_{1:m} \in S^m} \vec{w} \cdot \vec{\Phi}(x_{1:m}, s_{1:m}) \\
&= arg \max_{s_{1:m} \in S^m} \vec{w} \cdot \sum_{j=1}^{m} \vec{\phi}(x_{1:m}, j, s_{j-1}, s_j) \\
&= arg \max_{s_{1:m} \in S^m} \sum_{j=1}^{m} \vec{w} \cdot \vec{\phi}(x_{1:m}, j, s_{j-1}, s_j)
\end{aligned}
$$

$$(5)$$

Decoding with CRFs

- We have shown that finding the most likely sequence under the model is equivalent to finding the sequence that maximizes:

$$arg \max_{s_{1:m} \in \mathcal{S}^m} \sum_{j=1}^{m} \vec{w} \cdot \vec{\phi}(x_{1:m}, j, s_{j-1}, s_j)$$

- This problem has a clear intuition. Each transition from tag s_{j-1} to tag s_j has an associated score: $\vec{w} \cdot \vec{\phi}(x_{1:m}, j, s_{j-1}, s_j)$

- This score could be positive or negative.

- Intuitively, this score will be relatively high if the state transition is plausible, relatively low if this transition is implausible.

- The decoding problem is to find an entire sequence of states such that the sum of transition scores is maximized.

- We can again solve this problem using a variant of the Viterbi algorithm, in a very similar way to the decoding algorithm for HMMs or MEMMs.

Parameter Estimation in CRFs (training)

- For parameter estimation, we assume we have a set of n labeled examples, $\{(x^i_{1:m}, s^i_{1:m})\}^n_{i=1}$. Each $x^i_{1:m}$ is an input sequence $x^i_1, \ldots, x^i_m$ each $s^i_{1:m}$ is a tag sequence $s^i_1, \ldots, s^i_m$.

- We again set the negative log-likelihood (or cross-entropy) as the loss function L as optimize parameters using gradient descent.

- The main challenge here is that gradient calculations $\frac{\partial L}{\partial \vec{w}_{[k]}}$ involve summing over S^m (a very large set containing all possible tag sequences).

- This sum can be computed efficiently using the Forward-backward algorithm[4].

- This is another dynamic programming algorithm that is closely related to the Viterbi algorithm.

---

[4] http://www.cs.columbia.edu/~mcollins/fb.pdf

CRFs and MEMMs

- CRFs and MEMMS are discriminative sequence labeling models: they model the conditional probability directly via a parameterized log-linear multi-class function (softmax).
- HMMs, on the other hand, are generative models.
- In MEMM the normalization (denominator of the softmax) is local: it happens at each tag step (the sum runs over all possible tag values S).
- In CRFs the normalization is global: the sum runs over all possible tag sequences S^m.
- Training a MEMM is quite easy: just train a multi-class log-linear model for for a given word to the label. This classifier is used at each word step to predict the whole sequence.
- Training CRF is more complex. The objective is to maximize the log probability of the most likely sequence.

CRFs and MEMMs: the label bias problem

- MEMMs end up making up decision at each time step independently.
- This leads to a problem called label bias: in some tag space configurations, MEMMs essentially completely ignore important aspects of the context.
- Example: The right POS labeling of sentence "will to fight" (la voluntad de pelear) is "NN TO VB". [5]
- Here NN stands for "noun", TO stands for "infinitive to", and VB stands for "verb base form".
- Modals (MD) show up much more frequently at the start of the sentence than nouns do (e.g., questions).
- Hence, tag "MD" will receive a higher score than tag "NN" when x_0="will" : $P(s_1 = MD | s_0 = *, x_1 = \text{"will"}, ...) > P(s_1 = NN | s_{i-1} = *, x_1 = \text{"will"})$.
- But we know that MD + TO is very rare: "... can to eat", "... would to eat".

---

[5]Here we are using the PENN Treebank tagset:
https://www.eecis.udel.edu/~vijay/cis889/ie/pos-set.pdf

CRFs and MEMMs: the label bias problem

- The word "to" is relatively deterministic (almost always has tag TO) so it doesn't matter what tag precedes it.
- Because of the local normalization of MEMMs, $P(s_i = TO|s_{i-1}, x_1, \ldots, x_i = \text{"to"}, \ldots, x_n)$ will always be 1 when $x_i = \text{"to"}$ regardless of the value of s_{i-1} (MD or NN).
- That means our prediction for "to" can't help us disambiguate "will".
- We lose the information that MD + TO sequences rarely happen.
- As a consequence: a MEMMS would likely label the first word to "MD".
- CRF overcomes this issue by doing a global normalization: it considers the score of the whole sequence before normalizing to make it a probability distribution.

Links

- http://people.ischool.berkeley.edu/~dbamman/nlpF18/slides/ 11_memm_crf.pdf
- http://people.ischool.berkeley.edu/~dbamman/nlpF18/slides/ 12_neural_sequence_labeling.pdf
- https://www.depends-on-the-definition.com/ sequence-tagging-lstm-crf/
- https://www.quora.com/ What-are-the-pros-and-cons-of-these-three-sequence-models-MaxEnt-
- https: //people.cs.umass.edu/~mccallum/papers/crf-tutorial.pdf
- http://www.davidsbatista.net/blog/2017/11/13/Conditional_ Random_Fields

Questions?

Thanks for your Attention!

References I

Eisenstein, J. (2018).
Natural language processing.
Technical report, Georgia Tech.

Goodfellow, I., Bengio, Y., and Courville, A. (2016).
Deep learning.
MIT press.

Lafferty, J. D., McCallum, A., and Pereira, F. C. N. (2001).
Conditional random fields: Probabilistic models for segmenting and labeling sequence data.
In *Proceedings of the Eighteenth International Conference on Machine Learning*, ICML '01, pages 282–289, San Francisco, CA, USA. Morgan Kaufmann Publishers Inc.

McCallum, A., Freitag, D., and Pereira, F. C. (2000).
Maximum entropy markov models for information extraction and segmentation.
In *Icml*, pages 591–598.