

Natural Language Processing

Vector Space Model and Information Retrieval

Felipe Bravo-Marquez

March 22, 2020

Motivation

- How does a search engine such as Duckduckgo or Google retrieve relevant documents from a given query?
- How can a company process the claims left by its users on its Web portals?

These problems are studied in the following fields:

- *Information Retrieval*: science of searching for information in document collections.
- *Text Mining*: automatic extraction of knowledge from text.

Both of them are closely related to NLP! (the borders between these fields are unclear).

Tokens and Types

Tokenization: the task of splitting a sentence or document into pieces called *tokens*.

Additional transformations can be employed such as the removal of special characters (e.g., punctuation), lowercasing, etc. [Manning et al., 2008].

Example

Input: I like human languages and programming languages.

Tokens: [I] [like] [human] [languages] [and] [programming] [languages]

Types

- A *type* is a class of *token* containing a single sequence of characters.
- They are obtained by identifying unique tokens within the document.

Types for the previous sentence: [I] [like] [human] [languages] [and] [programming]

The token *languages* was repeated in the sentence.

Vocabulary Extraction

- A *term* is a normalized *type*.
- Normalization is the process of creating equivalence classes of different *types*. This will become clear in the following slides.
- The vocabulary V , is the set of terms (normalized unique tokens) within a collection of documents or corpus D .

Stopwords removal

- In order to reduce the size of the vocabulary and eliminate terms that do not provide much information, terms that occur with high frequency in the corpus are eliminated.
- These terms are called *stopwords* and include articles, pronouns, prepositions and conjunctions.

Example: [a, an, and, any, has, do, don't, did, the, on].¹

The removal of stopwords can be inconvenient in many NLP tasks!!

Example: I don't like pizza => pizza ("I", "don't", and "like" were removed)

¹Related concepts: function words, closed-class words.

Stemming

A term normalization process in which terms are transformed to their root in order to reduce the size of the vocabulary. It is carried by applying word reduction rules.

Example: Porter's Algorithm.

(F)	Rule	Example
	SSES → SS	caresses → caress
	IES → I	ponies → poni
	SS → SS	caress → caress
	S →	cats → cat

Example: d = I like human languages and programming languages \Rightarrow I like human languag and program languag²

The vocabulary of document d after removing stopwords and performing stemming:

termId	value
t1	human
t2	languag
t3	program

²http://9ol.es/porter_js_demo.html

Lemmatization

- Another term normalization strategy.
- It also transform words into their roots.
- It performs a morphological analysis using reference dictionaries (lookup tables) to create equivalence classes between *types*.
- For example, for the token *studies*, a stemming rule would return the term *studi*, while through lemmatization we would get the term *study*³.

³<https://blog.bitext.com/>

Zipf's law [1]

- The Zipf's law, proposed by *George Kingsley Zipf* in [Zipf, 1935], is an empirical law about the frequency of terms within a collection of documents (**corpus**).
- It states that the frequency f of a term in a corpus is inversely proportional to its r ranking in a sorted frequency table:

$$f = \frac{cf}{r^\beta} \quad (1)$$

- Where cf is a constant dependent on the collection and $\beta > 0$ is a decay factor.
- If $\beta = 1$, then f follows exactly Zipf's law, otherwise it follows a Zipf-like distribution.
- The law relates to the principle of minimum effort. We often use a few words to write ideas.
- The Zipf law is a type of power law distribution (long tail distributions)

Zipf's law [2]




Figure: Zipf's law

- If we plot a *log-log* graph, we obtain a straight line with slope $-\beta$.
- Listing the most frequent words of a corpus can be used to build a *stopwords* list.

Posting Lists and the Inverted Index

Let D be a collection of documents and V the vocabulary of all terms extracted from the collection:

- The posting list of a term is the list of all documents where the term appears at least once. Documents are identified by their ids.
- An inverted index is a dictionary-type data structure mapping terms $t_i \in V$ into their corresponding posting lists.

$$\langle term \rangle \rightarrow \langle docId \rangle^*$$




Figure: Inverted Index

Web Search Engines [1]

A search engine is an information retrieval system designed for searching information on the Web (solving information needs) [Manning et al., 2008]. Its basic components are:

- Crawler: a robot that navigates the Web according to a defined strategy. It usually starts by browsing a set of seed websites and continues to browse their hyperlinks.
- Indexer: in charge of maintaining an inverted index with the content of the pages traversed by the Crawler.
- Query processor: in charge of processing user queries and searching the index for the documents most relevant to a query.
- Ranking function: the function used by the query processor to rank documents indexed in the collection by relevance according to a query.
- User interface: receives the query as input and returns the documents ranked by relevancy.

Web Search Engines [2]

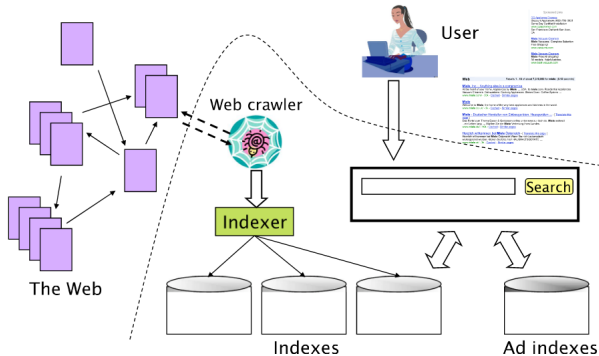


Figure: The various components of a web search engine [Manning et al., 2008].

The Vector Space Model

- In order to rank queries, or measure the similarity between two documents we need a similarity metric.
- Documents can be *represented* as vectors of terms, where each term is a vector dimension [Salton et al., 1975].
- Documents with different words and lengths will reside in the same vector space.
- These types of representations are called *Bag of Words*.
- In bag-of-words-representations the order of words and the linguistic structure of a sentence is lost.
- The value of each dimension is a weight that represents the relevance of the term t_i in the document d .

$$d_j \rightarrow \vec{d_j} = (w(t_1, d_j), \dots, w(t_{|V|}, d_j)) \quad (2)$$

- How can we model how informative is a term to a document?

Term Frequency - Inverted Document Frequency [1]

- Let $tf_{i,j}$ be the frequency of term t_i in document d_j .
- A term that occurs 10 times should provide more information than one that occurs once.
- What happens when we have documents that are much longer than the others?
- We can normalize by the maximum term frequency in the document.

$$ntf_{i,j} = \frac{tf_{i,j}}{\max_i(tf_{i,j})}$$

- Does a term that occurs in very few documents provide more or less information than one that occurs several times?
- For example, the document *The respected major of Pelotillehue*. The term *Pelotillehue* occurs in fewer documents than the term *major*, so it should be more descriptive.

Term Frequency - Inverted Document Frequency [2]

- Let N be the number of documents in the collection and n_i the number of documents containing term t_i , we define *idf* of t_i as follows:

$$idf_{t_i} = \log_{10}\left(\frac{N}{n_i}\right)$$

- A term that appears in all documents would have $idf = 0$ and one that appears in 10% of the documents would have $idf = 1$.
- The *tf-idf* scoring model combines *tf* and *idf* scores, resulting in the following weights w for a term in a document:

$$w(t_i, d_j) = tf_i \times \log_{10}\left(\frac{N}{n_i}\right)$$

- Search engine queries can also be modeled as vectors. However, queries have between 2 and 3 terms in average. To avoid having too many null dimensions, query vectors can be smoothed as follows:

$$w(t_i, d_j) = (0.5 + 0.5 \times tf_{i,j}) \log_{10}\left(\frac{N}{n_i}\right)$$

Similarity between Vectors

- Representing queries and documents as vectors allows calculating their similarity.
- One approach would be using the euclidean distance.
- The common approach is to calculate the cosine of the angle between the two vectors.
- If both documents are the same, the angle would be 0 and its cosine would be 1. On the other hand, if they are orthogonal the cosine is 0.
- The cosine similarity is calculated as follows:

$$\cos(\vec{d}_1, \vec{d}_2) = \frac{\vec{d}_1 \cdot \vec{d}_2}{|\vec{d}_1| \times |\vec{d}_2|} = \frac{\sum_{i=1}^{|V|} (w(t_i, d_1) \times w(t_i, d_2))}{\sqrt{\sum_{i=1}^{|V|} w(t_i, d_1)^2} \times \sqrt{\sum_{i=1}^{|V|} w(t_i, d_2)^2}}$$

- This is wrongly called *cosine distance*. It is actually a similarity metric.
- Notice that cosine similarity normalizes the vectors by its euclidean norm $||\vec{d}||_2$.

Cosine Similarity




Figure: Cosine Similarity.

Exercise

- Suppose we have 3 documents formed from the following sequences of terms:

$$d_1 \rightarrow t_4 t_3 t_1 t_4$$

$$d_2 \rightarrow t_5 t_4 t_2 t_3 t_5$$

$$d_3 \rightarrow t_2 t_1 t_4 t_4$$

- Build a term-document matrix of 5×3 dimensions using simple *tf-idf* weights (without normalization).
- We recommend you first build a list with the number of documents in which each term appears (useful for calculating *idf* scores)
- Then calculate the *idf* scores of each term.
- Fill up the cells of the matrix the *tf-idf* values.
- Which is the closest document to d_1 ?

Result

Table: tf-idf Matrix

	d1	d2	d3
t1	0.176	0.000	0.176
t2	0.000	0.176	0.176
t3	0.176	0.176	0.000
t4	0.000	0.000	0.000
t5	0.000	0.954	0.000

Document Clustering [1]

- How can we group documents that are similar with each other?
- Clustering is the process of grouping documents that are similar with each other.
- Each group of documents is called a *cluster*.
- In clustering we try to identify groups of documents in which the similarity between documents in the same cluster is maximized and the similarity of documents in different clusters is minimized.




Figure: Set of documents where the clusters can be clearly identified.

Document Clustering [2]

- Document clustering allows identifying topics in a corpus and reducing the search space in a search engine i.e., the inverted index is organized according to the clusters.
- K-means is a simple clustering algorithm that receives the number of clusters k as a parameter.
- The algorithm relies on the idea of *centroid*, which is the average vector of documents belonging to the same cluster.
- Let S be a set of 2-dimensional vectors $\{3, 6\}$, $\{1, 2\}$, $\{5, 1\}$, the centroid of S is $\{(3 + 1 + 5)/3, (6 + 2 + 1)/3\} = \{3, 3\}$.

K-Means

1. We start with k random centroids.
2. We calculate the similarity between each document and each centroid.
3. We assign each document to its closest centroid forming a cluster.
4. The centroids are recalculated according to the documents assigned to them.
5. This process is repeated until convergence.

K-means

```
K-MEANS( $\{\vec{x}_1, \dots, \vec{x}_N\}, K$ )
1   $(\vec{s}_1, \vec{s}_2, \dots, \vec{s}_K) \leftarrow \text{SELECTRANDOMSEEDS}(\{\vec{x}_1, \dots, \vec{x}_N\}, K)$ 
2  for  $k \leftarrow 1$  to  $K$ 
3  do  $\vec{\mu}_k \leftarrow \vec{s}_k$ 
4  while stopping criterion has not been met
5  do for  $k \leftarrow 1$  to  $K$ 
6      do  $\omega_k \leftarrow \{\}$ 
7      for  $n \leftarrow 1$  to  $N$ 
8          do  $j \leftarrow \arg \min_j |\vec{\mu}_j - \vec{x}_n|$ 
9               $\omega_j \leftarrow \omega_j \cup \{\vec{x}_n\}$  (reassignment of vectors)
10     for  $k \leftarrow 1$  to  $K$ 
11         do  $\vec{\mu}_k \leftarrow \frac{1}{|\omega_k|} \sum_{\vec{x} \in \omega_k} \vec{x}$  (recomputation of centroids)
12 return  $\{\vec{\mu}_1, \dots, \vec{\mu}_K\}$ 
```

Figure: K-means algorithm

Conclusions and Additional Concepts

- Representing documents as vectors is essential for calculating similarities between document pairs.
- Bag of words vectors lack linguistic structure.
- Bag of words vectors are high-dimensional and sparse.
- Word n-grams can help capturing multi word-expressions (e.g., New York => new_york)
- Modern information retrieval systems go beyond vector similarity (PageRank, Relevance Feedback, Query log mining, Google Knowledge Graph, Machine Learning).
- Information retrieval and text mining are less concerned with linguistic structure, and more interested in producing fast and scalable algorithms [Eisenstein, 2018].

References I



Eisenstein, J. (2018).
Natural language processing.
Technical report, Georgia Tech.



Manning, C. D., Raghavan, P., and Schütze, H. (2008).
Introduction to Information Retrieval.
Cambridge University Press, New York, NY, USA.



Salton, G., Wong, A., and Yang, C.-S. (1975).
A vector space model for automatic indexing.
Communications of the ACM, 18(11):613–620.



Zipf, G. K. (1935).
The Psychobiology of Language.
Houghton-Mifflin, New York, NY, USA.